



# contents

---

<i>foreword</i>	<i>xiii</i>
<i>letter to the HBase community</i>	<i>xv</i>
<i>preface</i>	<i>xvii</i>
<i>acknowledgments</i>	<i>xix</i>
<i>about this book</i>	<i>xxi</i>
<i>about the authors</i>	<i>xxv</i>
<i>about the cover illustration</i>	<i>xxvi</i>

## PART 1 HBASE FUNDAMENTALS . . . . .1

### 1 Introducing HBase 3

1.1	Data-management systems: a crash course	5			
	<i>Hello, Big Data</i>	6 ▪ <i>Data innovation</i>	7 ▪ <i>The rise of HBase</i>	8	
1.2	HBase use cases and success stories	8			
	<i>The canonical web-search problem: the reason for Bigtable's invention</i>	9 ▪ <i>Capturing incremental data</i>	10 ▪ <i>Content serving</i>	13 ▪ <i>Information exchange</i>	14
1.3	Hello HBase	15			
	<i>Quick install</i>	16 ▪ <i>Interacting with the HBase shell</i>	18		
	<i>Storing data</i>	18			
1.4	Summary	20			

## 2 *Getting started* 21

### 2.1 Starting from scratch 22

*Create a table* 22 ▪ *Examine table schema* 23 ▪ *Establish a connection* 24 ▪ *Connection management* 24

### 2.2 Data manipulation 25

*Storing data* 25 ▪ *Modifying data* 26 ▪ *Under the hood: the HBase write path* 26 ▪ *Reading data* 28 ▪ *Under the hood: the HBase read path* 29 ▪ *Deleting data* 30 ▪ *Compactions: HBase housekeeping* 30 ▪ *Versioned data* 31 ▪ *Data model recap* 32

### 2.3 Data coordinates 33

### 2.4 Putting it all together 35

### 2.5 Data models 39

*Logical model: sorted map of maps* 39 ▪ *Physical model: column family oriented* 41

### 2.6 Table scans 42

*Designing tables for scans* 43 ▪ *Executing a scan* 45 ▪ *Scanner caching* 45 ▪ *Applying filters* 46

### 2.7 Atomic operations 47

### 2.8 ACID semantics 48

### 2.9 Summary 48

## 3 *Distributed HBase, HDFS, and MapReduce* 51

### 3.1 A case for MapReduce 52

*Latency vs. throughput* 52 ▪ *Serial execution has limited throughput* 53 ▪ *Improved throughput with parallel execution* 53 ▪ *MapReduce: maximum throughput with distributed parallelism* 55

### 3.2 An overview of Hadoop MapReduce 56

*MapReduce data flow explained* 57 ▪ *MapReduce under the hood* 61

### 3.3 HBase in distributed mode 62

*Splitting and distributing big tables* 62 ▪ *How do I find my region?* 64 ▪ *How do I find the -ROOT- table?* 65

### 3.4 HBase and MapReduce 68

*HBase as a source* 68 ▪ *HBase as a sink* 70 ▪ *HBase as a shared resource* 71

- 3.5 Putting it all together 75
  - Writing a MapReduce application* 76 ▪ *Running a MapReduce application* 77
- 3.6 Availability and reliability at scale 78
  - HDFS as the underlying storage* 79
- 3.7 Summary 81

## PART 2 ADVANCED CONCEPTS.....83

### 4 HBase table design 85

- 4.1 How to approach schema design 86
  - Modeling for the questions* 86 ▪ *Defining requirements: more work up front always pays* 89 ▪ *Modeling for even distribution of data and load* 92 ▪ *Targeted data access* 98
- 4.2 De-normalization is the word in HBase land 100
- 4.3 Heterogeneous data in the same table 102
- 4.4 Rowkey design strategies 103
- 4.5 I/O considerations 104
  - Optimized for writes* 104 ▪ *Optimized for reads* 106
  - Cardinality and rowkey structure* 107
- 4.6 From relational to non-relational 108
  - Some basic concepts* 109 ▪ *Nested entities* 110 ▪ *Some things don't map* 112
- 4.7 Advanced column family configurations 113
  - Configurable block size* 113 ▪ *Block cache* 114 ▪ *Aggressive caching* 114 ▪ *Bloom filters* 114 ▪ *TTL* 115
  - Compression* 115 ▪ *Cell versioning* 116
- 4.8 Filtering data 117
  - Implementing a filter* 119 ▪ *Prebundled filters* 121
- 4.9 Summary 124

### 5 Extending HBase with coprocessors 126

- 5.1 The two kinds of coprocessors 127
  - Observer coprocessors* 127 ▪ *Endpoint Coprocessors* 130
- 5.2 Implementing an observer 131
  - Modifying the schema* 131 ▪ *Starting with the Base* 132
  - Installing your observer* 135 ▪ *Other installation options* 137

- 5.3 Implementing an endpoint 137
  - Defining an interface for the endpoint* 138 ▪ *Implementing the endpoint server* 138 ▪ *Implement the endpoint client* 140
  - Deploying the endpoint server* 142 ▪ *Try it!* 142
- 5.4 Summary 142

## 6 *Alternative HBase clients* 143

- 6.1 Scripting the HBase shell from UNIX 144
  - Preparing the HBase shell* 145 ▪ *Script table schema from the UNIX shell* 145
- 6.2 Programming the HBase shell using JRuby 147
  - Preparing the HBase shell* 147 ▪ *Interacting with the TwitBase users table* 148
- 6.3 HBase over REST 150
  - Launching the HBase REST service* 151 ▪ *Interacting with the TwitBase users table* 153
- 6.4 Using the HBase Thrift gateway from Python 156
  - Generating the HBase Thrift client library for Python* 157
  - Launching the HBase Thrift service* 159 ▪ *Scanning the TwitBase users table* 159
- 6.5 Asynchbase: an alternative Java HBase client 162
  - Creating an asynchbase project* 163 ▪ *Changing TwitBase passwords* 165 ▪ *Try it out* 176
- 6.6 Summary 177

## PART 3 *EXAMPLE APPLICATIONS*.....179

### 7 *HBase by example: OpenTSDB* 181

- 7.1 An overview of OpenTSDB 182
  - Challenge: infrastructure monitoring* 183 ▪ *Data: time series* 184
  - Storage: HBase* 185
- 7.2 Designing an HBase application 186
  - Schema design* 187 ▪ *Application architecture* 190
- 7.3 Implementing an HBase application 194
  - Storing data* 194 ▪ *Querying data* 199
- 7.4 Summary 202

## 8 *Scaling GIS on HBase* 203

- 8.1 Working with geographic data 203
- 8.2 Designing a spatial index 206
  - Starting with a compound rowkey* 208 ▪ *Introducing the geohash* 209 ▪ *Understand the geohash* 211 ▪ *Using the geohash as a spatially aware rowkey* 212
- 8.3 Implementing the nearest-neighbors query 216
- 8.4 Pushing work server-side 222
  - Creating a geohash scan from a query polygon* 224 ▪ *Within query take 1: client side* 228 ▪ *Within query take 2: WithinFilter* 231
- 8.5 Summary 235

## PART 4 OPERATIONALIZING HBASE .....237

## 9 *Deploying HBase* 239

- 9.1 Planning your cluster 240
  - Prototype cluster* 241 ▪ *Small production cluster (10–20 servers)* 242
  - Medium production cluster (up to ~50 servers)* 243 ▪ *Large production cluster (>~50 servers)* 243 ▪ *Hadoop Master nodes* 243 ▪ *HBase Master* 244 ▪ *Hadoop DataNodes and HBase RegionServers* 245
  - ZooKeeper(s)* 246 ▪ *What about the cloud?* 246
- 9.2 Deploying software 248
  - Whirr: deploying in the cloud* 249
- 9.3 Distributions 250
  - Using the stock Apache distribution* 251 ▪ *Using Cloudera's CDH distribution* 252
- 9.4 Configuration 253
  - HBase configurations* 253 ▪ *Hadoop configuration parameters relevant to HBase* 260 ▪ *Operating system configurations* 261
- 9.5 Managing the daemons 261
- 9.6 Summary 263

## 10 *Operations* 264

- 10.1 Monitoring your cluster 265
  - How HBase exposes metrics* 266 ▪ *Collecting and graphing the metrics* 266 ▪ *The metrics HBase exposes* 268 ▪ *Application-side monitoring* 272